

## Bayesian regularization — Project

The project consists in implementing one of the two classification methods seen in the courses. You have the choice between quadratic discriminant analysis and binary logistic regression. You must implement the regularized version of the chosen method, and then apply it to the datasets provided.

You will have to provide a report (L<sup>A</sup>T<sub>E</sub>X recommended, 10 pages maximum) which *briefly* recalls the main equations behind each classification method; which describes your implementation; and which provides a presentation and some comments on the results obtained.

You will also provide the Python code with your functions and the script with your training/test procedure.

The projects are due for Dec. 20<sup>th</sup>, 2019.

### 1 Quadratic Discriminant Analysis (QDA)

**Prior choice** Recall that the Gaussian-inverse-Wishart prior is

$$\mu_k \sim \mathcal{N}(\mu_{kp}, \Sigma_k / \kappa_{kp}), \quad \Sigma_k \sim IW(\nu_{kp}, \Lambda_{kp});$$

the meta-parameters of these distributions should satisfy  $\kappa_{kp} > 0$  and  $\nu_{kp} > d - 1$  (with  $d$  standing for the dimension of the space). The covariance matrix for  $\mu_k$  will be finite provided  $\nu_{kp} > d + 1$ .

The priors can be interpreted as information provided by “virtual data”; they should make it possible to make model estimation more robust without involving any information regarding the actual distribution of the data (remember that while with simulated data the actual distribution is known, this will not be the case for real datasets).

You have to imagine a way to provide meaningful estimates using only the observed training data at your disposal. You can test various values for the shrinkage parameter and degrees of freedom and analyze the influence on the results.

**Code** The function for training QDA should take the following inputs :

- the  $n \times d$  matrix **Xapp** of instances  $\mathbf{x}_i \in \mathbb{R}^d$  (instances are represented row-wise),
- the  $n \times K$  matrix **zapp** of (row-wise) binary class indicators;
- the  $K \times p$  matrix **mprior** of expectations  $\mu_{kp}$  for the Gaussian prior on  $\mu_k$ ,
- the shrinkage parameter **df\_exp** of the Gaussian prior ( $\kappa_{kp}$ ),
- the  $p \times p \times K$  matrix **Sprior** of covariance matrices  $\Lambda_{kp}^{-1}$  for the inverse-Wishart prior on  $\Sigma_k$ ,
- the number **df\_cov** of degrees of freedom of the inverse-Wishart prior ( $\nu_{kp} \geq d - 1$ ).

It should provide :

- the  $1 \times K$  vector **pi** of class proportions,
- the  $K \times p$  matrix **mu** of estimated expectations,
- the  $p \times p \times K$  matrix **Sig** of estimated covariance matrices.

The function for evaluating test data should take the following inputs :

- the  $n \times p$  matrix **Xtst** of test instances  $\mathbf{x}_i \in \mathbb{R}^p$  (instances are represented row-wise),
- the  $1 \times K$  vector **pi** of class proportions, the  $K \times p$  matrix **mu** of estimated expectations, and the  $p \times p \times K$  matrix **Sig** of estimated covariance matrices;

and it should provide

- the  $n \times K$  matrix **prob** of estimated class posterior probabilities,
- the  $n \times 1$  vector **pred** of corresponding decisions.

**Tests** Process the data by varying the amount of available training data. Compare the results obtained without and with priors. You will consider the following datasets : Optdigits, Pageblocks, Satimage, Segment.

## 2 Logistic regression

**Prior choice** You will test both priors on the vector parameter  $\beta$  mentioned during the course : the Gaussian prior and the Laplacean prior.

**Code** The logistic regression code provided by the function **logfit.m** takes the following inputs :

- the  $n \times d$  matrix **X** of instances  $\mathbf{x}_i \in \mathbb{R}^d$  (instances are represented row-wise),
- the  $n \times K$  matrix **y** of binary class indicators  $z_{ik}$  (row-wise, possibly imprecise : whenever the constraint  $\sum_k z_{ik} = 1$  is relaxed, they provide the sets of possible classes for the instances) ;
- a flag parameter **intercept** indicating whether an intercept should be used in the model, or not ;
- the scalar **lambda** used for regularization,
- the precision scalar **precision** used to check for convergence,
- the initial parameter **betaini** standing for the initial vector of parameters  $\beta^{(0)}$ .

It should provide :

- the  $p \times 1$  (no intercept) or  $(p + 1) \times 1$  (intercept) vector **beta** which corresponds to the maximum-likelihood estimate  $\hat{\beta}$  for  $\beta$  (or maximum a posteriori if  $\lambda > 0$ ),
- the log-likelihood **logL** corresponding to  $\log L(\hat{\beta}; \dots)$  (if  $\lambda = 0$ ) or  $\pi_{\beta}(\hat{\beta}; \dots)$  (if  $\lambda > 0$ ).

The test function should take as input

- the  $n \times p$  matrix **Xtst** of test instances  $\mathbf{x}_i \in \mathbb{R}^p$  (instances are represented row-wise),
- the  $p \times 1$  vector (or  $(p + 1) \times 1$  vector if intercept) **beta** of estimated parameters ;

and it should provide

- the  $n \times K$  matrix **prob** of estimated class posterior probabilities,
- the  $n \times 1$  vector **pred** of corresponding decisions.

### 2.1 Tests

Process the data by varying the amount of available training data. Compare the results obtained without and with priors. You will consider the following datasets : Breastcancer, Ionosphere, Sonar, Spambase.